

Balanced Search Trees

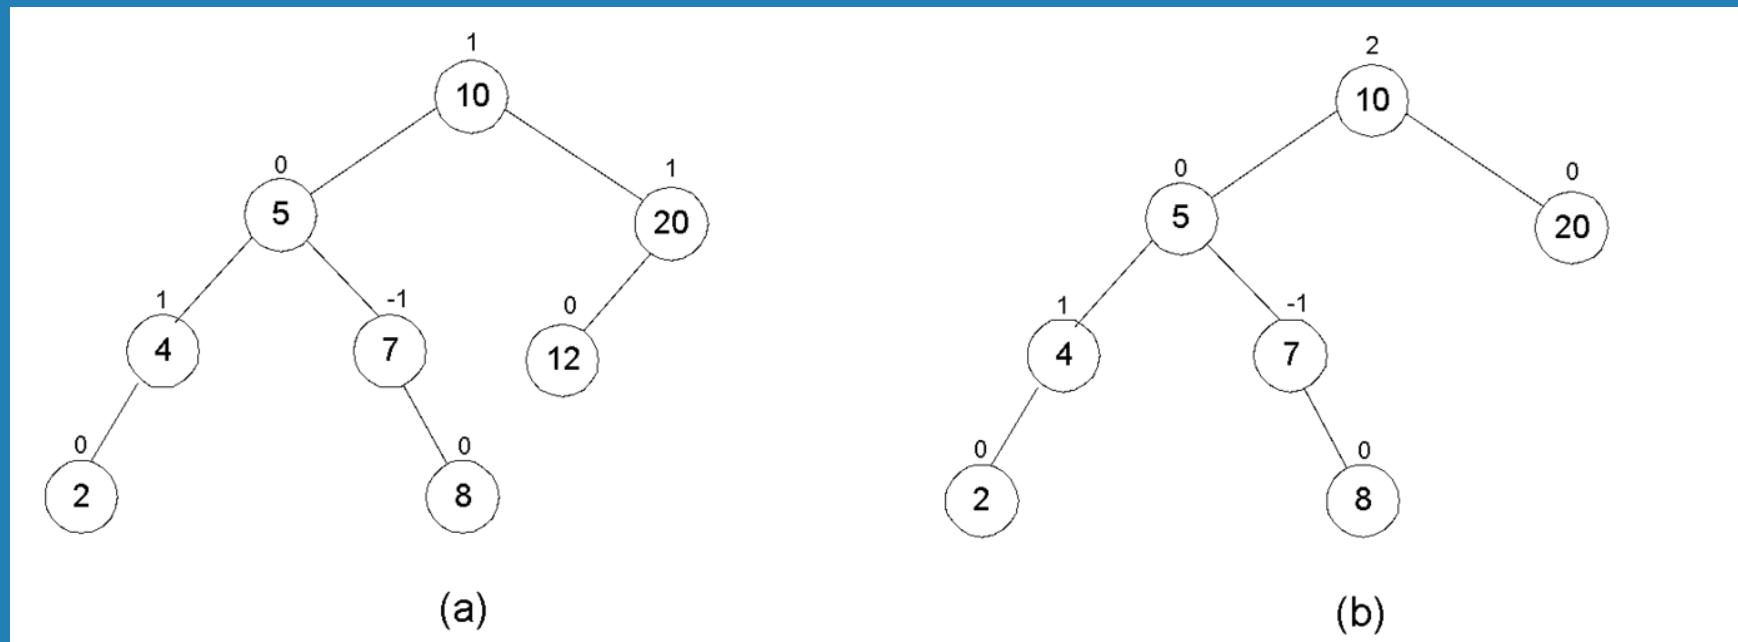


Attractiveness of *binary search tree* is marred by the bad (linear) worst-case efficiency. Two ideas to overcome it are:

- to rebalance binary search tree when a new insertion makes the tree “too unbalanced”
 - *AVL trees*
 - *red-black trees*
- to allow more than one key and two children
 - *2-3 trees*
 - *2-3-4 trees*
 - *B-trees*

Balanced trees: AVL trees

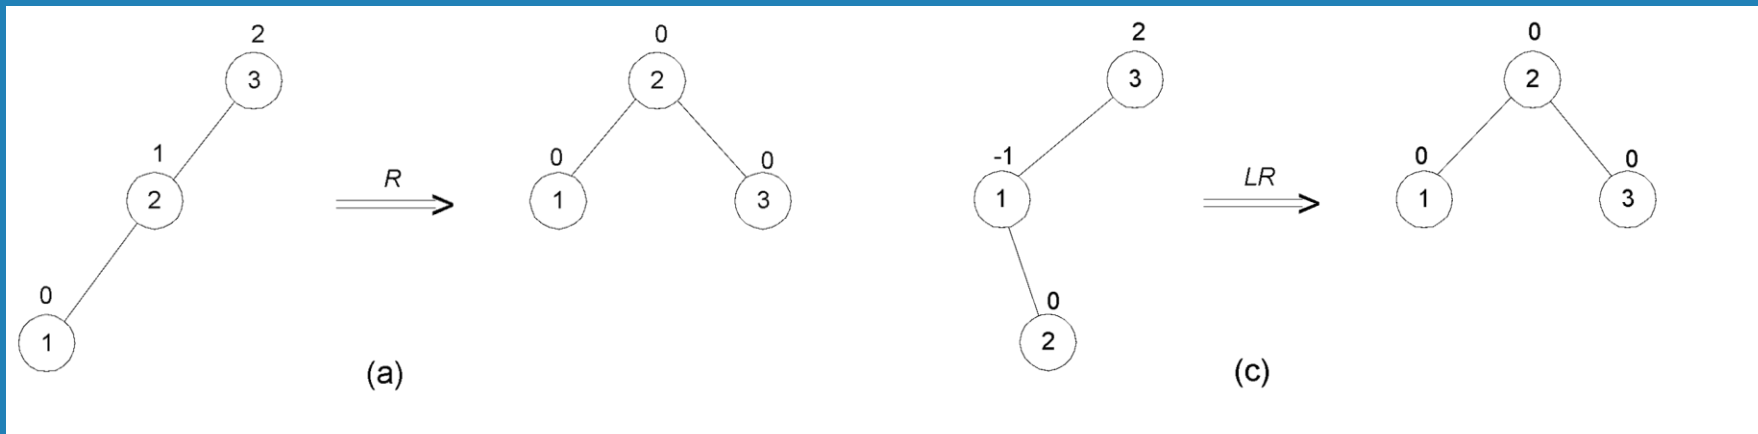
Definition An *AVL tree* is a binary search tree in which, for every node, the difference between the heights of its left and right subtrees, called the *balance factor*, is at most 1 (with the height of an empty tree defined as -1)



Tree (a) is an AVL tree; tree (b) is not an AVL tree

Rotations


If a key insertion violates the balance requirement at some node, the subtree rooted at that node is transformed via one of the four *rotations*. (The rotation is always performed for a subtree rooted at an “unbalanced” node closest to the new leaf.)




Single R -rotation

Double LR -rotation

General case: Single R-rotation



General case: Double LR-rotation



AVL tree construction - an example

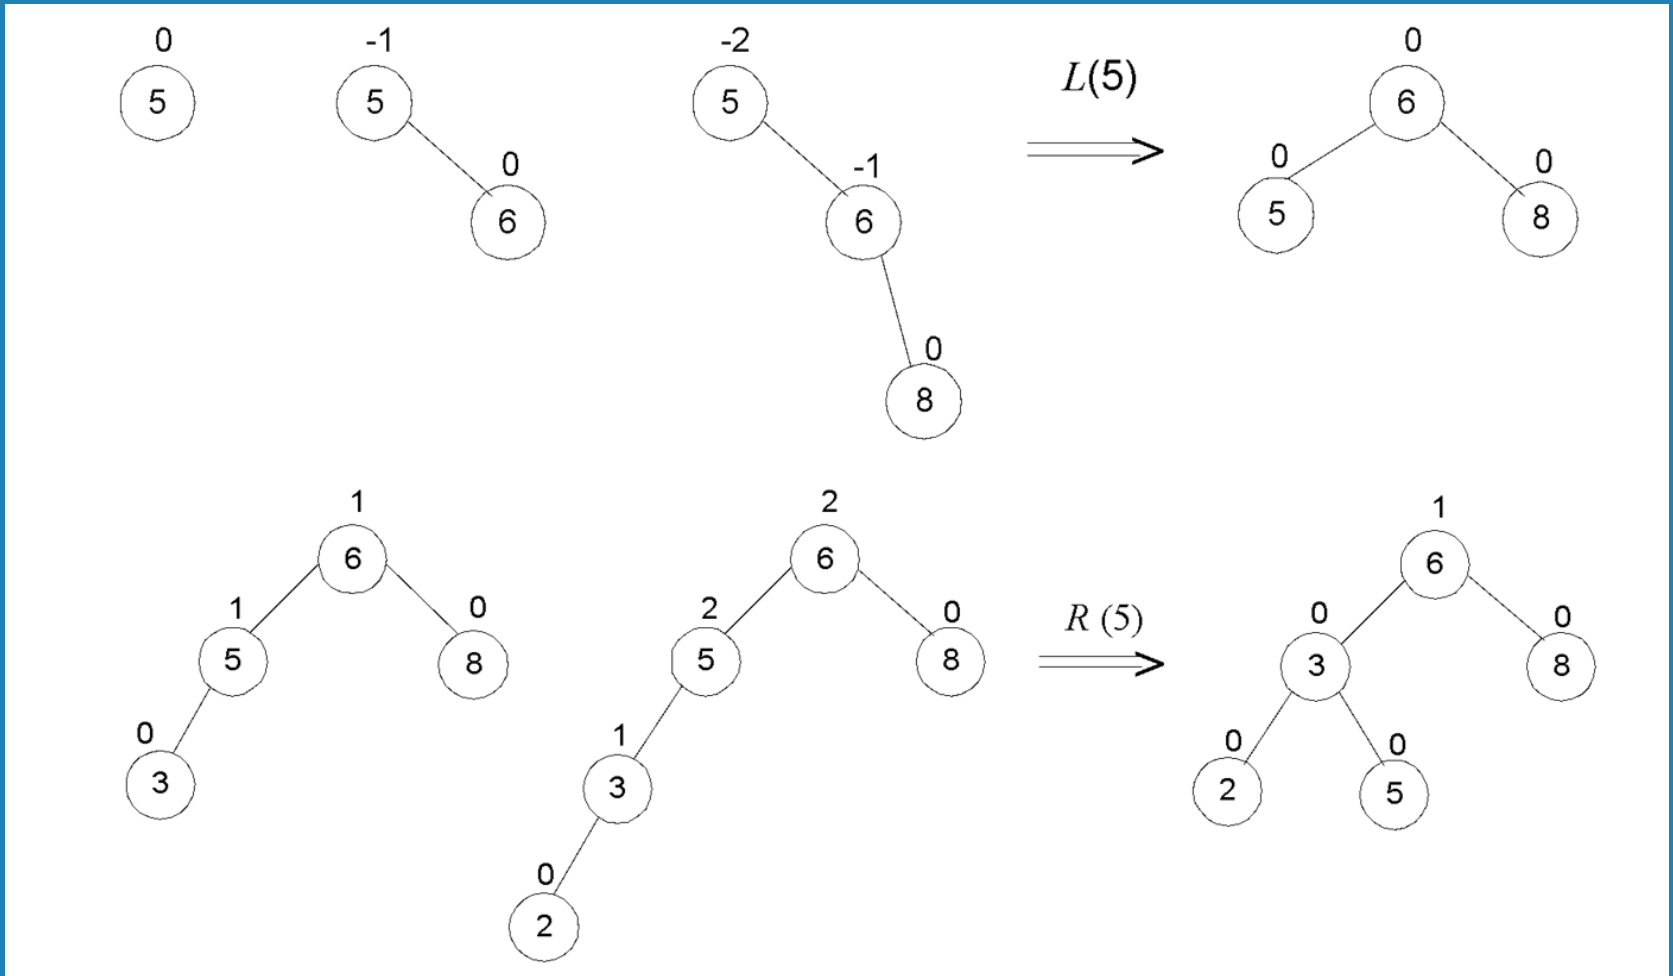


Construct an AVL tree for the list 5, 6, 8, 3, 2, 4, 7






AVL tree construction - an example

Construct an AVL tree for the list 5, 6, 8, 3, 2, 4, 7



AVL tree construction - an example (cont.)





Analysis of AVL trees



- $h \leq 1.4404 \log_2 (n + 2) - 1.3277$ $N(h-1) + N(h-2) \leq N(h)$
average height: $1.01 \log_2 n + 0.1$ for large n (found empirically)
- Search and insertion are $O(\log n)$
- Deletion is more complicated but is also $O(\log n)$
- Disadvantages:
 - frequent rotations
 - complexity
- A similar idea: *red-black trees* (height of subtrees is allowed to differ by up to a factor of 2)